

RANK-OPTIMAL DIRECTIONAL TAYLOR JETS FOR HIGH-ORDER DERIVATIVES

AUTHOR ONE* AND AUTHOR TWO†

Abstract. A prescribed monomial derivative is one term in a high-order residual in several variables, as in a polyharmonic operator. Nested automatic differentiation computes it exactly, but each additional coordinate derivative deepens the computational graph. Finite differences and randomized estimators avoid that nesting only by introducing a step size or a sampling variance. Taylor-mode automatic differentiation removes the nesting by propagating truncated Taylor coefficients along one direction at a time. The cost is then the number of directions needed for that derivative. We construct a directional Taylor-jet schedule whose length equals the Waring rank of the associated monomial, which is known in closed form over the complex numbers, and we realize it by a roots-of-unity formula together with a batched Taylor jet and a closed-form adjoint. The directional identity is algebraically exact and admits complex directions. Compared with nested automatic differentiation on the same collocation residual in double precision, the two methods agree for derivative values and for parameter gradients, and the schedule reduces the time and the memory of one evaluation as the derivative order grows.

Key words. high-order derivatives, Taylor-mode automatic differentiation, Waring decomposition, apolarity

AMS subject classifications. 65D25, 14N07, 35J40

1. Introduction. High-order residuals in several variables are assembled from monomial derivatives. Each monomial is evaluated repeatedly, at many points, on one composition that can be probed along arbitrary directions, and the cost of those evaluations grows with the order of the derivative and with the depth of the composition.

The exact methods are classical, and each pays for exactness differently. Symbolic differentiation produces an expression that grows rapidly with order. Finite differences are insensitive to the internal structure of the composition, but they introduce a step size, and both truncation and cancellation become more severe as the order rises [3]. Automatic differentiation differentiates the evaluation procedure itself [5]. Reverse mode yields a gradient in one backward pass; a higher-order derivative is typically obtained by repeating that process, and each repetition differentiates the graph built by the previous one, so depth and storage grow with the order. Taylor mode avoids the nesting: one forward sweep along a single direction returns the directional derivatives up to a chosen order [1, 5]. Its cost is therefore measured in directional sweeps, and the number of directions is the quantity to minimize.

Randomized estimators replace a derivative by an expectation and scale to high dimension at the price of variance [6–8, 14]. Operator-specific recurrences, such as the forward Laplacian, keep exactness but are tied to the operator they were designed for [4, 12]. The problem treated here is the evaluation of one prescribed monomial derivative, computed exactly, of a composition that can be evaluated along arbitrary directions.

Taylor mode computes the directional derivative it is given. It does not say how many directions a prescribed monomial needs. Elementary reconstructions exist, obtained by evaluating along combinations of coordinate axes, but a formula written by hand does not certify that a shorter one is impossible. The question is therefore the

*Institution One (author1@example.com).

†Institution Two (author2@example.com).

minimum: for one multi-index, of one order, what is the smallest number of directional Taylor evaluations that recover it exactly, and which directions attain that bound?

The directional Taylor coefficient of a fixed order is a homogeneous polynomial in the direction. Extracting the target derivative is extracting one monomial coefficient of that polynomial, so a directional formula is a decomposition of a monomial into powers of linear forms. The length of a shortest formula is the Waring rank of that monomial, which the algebraic-geometry literature has evaluated in closed form over the complex numbers. A roots-of-unity construction attains the rank. Because the rank is a minimum, no shorter linear combination of order- p directional Taylor coefficients exists.

We identify that shortest formula with the monomial Waring rank, and we give the roots-of-unity schedule that attains it. The schedule is realized as a batched Taylor jet through an entire activation, with a closed-form adjoint so that the result remains differentiable in the parameters of the composition. The experiments compare this method with nested automatic differentiation on otherwise identical collocation residuals: the two agree in double precision, and the reduction in the number of directions appears as a reduction in time and memory.

The remainder of the paper is organized as follows. Section 2 states the evaluation problem and recalls Taylor-mode directional differentiation. Section 3 identifies the shortest directional formula with a monomial Waring decomposition, proves the main theorem, and records the algorithm. Section 4 checks the implementation against nested automatic differentiation on collocation residuals and reports the cost of one evaluation. Section 5 concludes. The roots-of-unity filter, the explicit schedules used in the experiments, and the apolar form of the rank bound are collected in Appendix A.

2. Preliminaries.

2.1. The evaluation problem. Let $u_\theta: \mathbb{R}^d \rightarrow \mathbb{K}$ with $\mathbb{K} \in \{\mathbb{R}, \mathbb{C}\}$ be a parameterized composition, smooth in its inputs. Let $\nu = (\nu_1, \dots, \nu_d) \in \mathbb{Z}_{\geq 0}^d$ be a multi-index of order $p = |\nu|$, and write $\partial^\nu := \partial_1^{\nu_1} \cdots \partial_d^{\nu_d}$ for the associated derivative. The task is the evaluation of $\partial^\nu u_\theta$ at a given point. The value must be algebraically exact, and it must remain differentiable in θ . The multi-index is fixed and prescribed. The derivative is evaluated at many points at once and re-evaluated at every step, so the cost that matters is one evaluation of one derivative, amortized over a batch.

Section 4 measures that cost on collocation residuals of a boundary-value problem [10, 13]. Let $\Omega \subset \mathbb{R}^d$ be a bounded domain with boundary $\Gamma := \partial\Omega$, and consider

$$\mathcal{P}u = f \quad \text{in } \Omega, \quad \mathcal{B}u = g \quad \text{on } \Gamma, \quad (2.1)$$

with $\mathcal{P}$ a differential operator of order p and $\mathcal{B}$ the boundary conditions. Given interior points $(\mathbf{x}_{r,i})_{i=1}^{n_r} \subset \Omega$ and boundary points $(\mathbf{x}_{b,i})_{i=1}^{n_b} \subset \Gamma$, a collocation fit of (2.1) minimizes

$$\mathcal{L}(\theta) := \frac{\lambda_r}{n_r} \sum_{i=1}^{n_r} |\mathcal{P}u_\theta(\mathbf{x}_{r,i}) - f(\mathbf{x}_{r,i})|^2 + \frac{\lambda_b}{n_b} \sum_{i=1}^{n_b} |\mathcal{B}u_\theta(\mathbf{x}_{b,i}) - g(\mathbf{x}_{b,i})|^2. \quad (2.2)$$

Expanding $\mathcal{P}$ into monomials turns one residual evaluation in (2.2) into several derivatives $\partial^\alpha u_\theta$ of the same composition on the same batch. The operator Δ^m expands into $\binom{d+m-1}{m}$ distinct monomials: the pure terms $\partial_{x_i}^{2m}$ and, for $m \geq 2$, mixed terms such as $\partial_{x_i}^{2k} \partial_{x_j}^{2m-2k}$.

2.2. Taylor-mode evaluation of directional derivatives. Let $u: \mathbb{R}^d \rightarrow \mathbb{K}$ with $\mathbb{K} \in \{\mathbb{R}, \mathbb{C}\}$ be sufficiently smooth near a fixed point $\mathbf{x}$. For a direction $\mathbf{v} \in \mathbb{R}^d$, the directional Taylor coefficient of order p at $\mathbf{x}$ along $\mathbf{v}$ is

$$T_p(\mathbf{x}; \mathbf{v}) := \frac{1}{p!} \frac{d^p}{d\tau^p} u(\mathbf{x} + \tau \mathbf{v}) \Big|_{\tau=0} = \frac{1}{p!} D^p u(\mathbf{x})[\mathbf{v}, \dots, \mathbf{v}]. \quad (2.3)$$

Taylor-mode automatic differentiation evaluates (2.3) in one forward pass [1, 5]. Each intermediate activation $y \in \mathbb{K}^B$ over a batch of B points, with $\mathbb{K} \in \{\mathbb{R}, \mathbb{C}\}$, is replaced by the length- $(p+1)$ tuple $(y_0, \dots, y_p)$ of its truncated Taylor coefficients $y_k = (1/k!) d^k y / d\tau^k|_{\tau=0}$. For a multilayer perceptron $u = \phi_L \circ A_L \circ \dots \circ \phi_1 \circ A_1$ with affine maps A_ℓ and entrywise activations ϕ_ℓ , an affine layer $y = Wx + b$ propagates each order independently,

$$y_0 = Wx_0 + b, \quad y_k = Wx_k \quad (k \geq 1), \quad (2.4)$$

and each nonlinearity is propagated together with one auxiliary jet that closes its derivative recurrence. For $\tanh$ the auxiliary jet is $z = 1 - y^2$, initialized by $y_0 = \tanh(x_0)$ and $z_0 = 1 - y_0^2$, and

$$y_k = \frac{1}{k} \sum_{j=0}^{k-1} (k-j) z_j x_{k-j}, \quad z_k = - \sum_{j=0}^k y_j y_{k-j}, \quad k \geq 1. \quad (2.5)$$

For $\sinh$ the auxiliary jet is $z = \cosh(x)$, initialized by $y_0 = \sinh(x_0)$ and $z_0 = \cosh(x_0)$, and

$$y_k = \frac{1}{k} \sum_{j=0}^{k-1} (k-j) z_j x_{k-j}, \quad z_k = \frac{1}{k} \sum_{j=0}^{k-1} (k-j) y_j x_{k-j}, \quad k \geq 1. \quad (2.6)$$

The recurrence (2.6) has the same coupled form as the one for $\sin$ and $\cos$, except that the $\cosh$ update carries no sign flip. The implementation also contains the recurrences for $\sin$ and for the logistic sigmoid, which the verification tests of Section 4 use. Every coefficient y_k is an ordinary tensor, and the propagation stores $p+1$ coefficient tensors per activation without nesting input-gradient graphs.

3. Rank-optimal directional schedules. The argument of this section is short. Its four steps are as follows. First, the directional Taylor coefficients of order p , read as a function of the direction rather than of the point, are the coefficients of a homogeneous polynomial of degree p , and the target derivative is one of them. Second, a formula that reconstructs the derivative from finitely many directional evaluations is therefore a linear identity among such coefficients, which reduces the design problem to a finite system of moment conditions. Third, pairing each direction with the polynomial variable turns that system into a single polynomial identity: the monomial belonging to the target derivative, written as a linear combination of p th powers of linear forms. Fourth, the least number of terms in such a combination is by definition the Waring rank of that monomial, a quantity known in closed form, and one set of minimizing linear forms consists of roots of unity. The four steps together replace a search for a reconstruction formula by a rank computation, and because the outcome is a minimum, it certifies that no shorter linear combination of those coefficients exists.

We introduce some notation used throughout this section. Let $\mathbb{N}_0 := \{0, 1, 2, \dots\}$, and let $\mathbb{C}^d$ be the d -dimensional complex space. The *active coordinates* of ν are the indices k with $\nu_k > 0$; we relabel them as $i_0, \dots, i_n$ and set $a_j := \nu_{i_j}$, ordered so that $1 \leq a_0 \leq \dots \leq a_n$, with ties broken by the original coordinate index. The ordering singles out a coordinate i_0 of least exponent, which is the coordinate the rank formula treats separately. We write $a := (a_0, \dots, a_n)$, and we record for later use that

$$p = \sum_{j=0}^n a_j, \quad \nu! = \prod_{k=1}^d \nu_k! = \prod_{j=0}^n a_j!, \quad m_j := a_j + 1. \quad (3.1)$$

The *active subspace* is $V := \text{span}\{e_{i_0}, \dots, e_{i_n}\} \subset \mathbb{R}^d$, and $V_{\mathbb{C}} := V \otimes_{\mathbb{R}} \mathbb{C}$ is its complexification. Since ∂^ν is insensitive to coordinates outside the active set, all directional probes below may be taken in V or in $V_{\mathbb{C}}$. For $\mathbb{K} \in \{\mathbb{R}, \mathbb{C}\}$, the space of homogeneous polynomials of degree p in $\mathbf{z} = (z_0, \dots, z_n)$ is denoted $\mathbb{K}[\mathbf{z}]_p$, and $[\mathbf{z}^a]F$ denotes the coefficient of $\mathbf{z}^a = z_0^{a_0} \dots z_n^{a_n}$ in F . For $m \geq 1$, let $U_m := \{\zeta \in \mathbb{C} : \zeta^m = 1\}$, and let $\mathbf{1}_A$ be the indicator of a condition A .

The minimizing schedule of this section is in general complex, so we fix the meaning of a complex direction once. For $\mathbf{v} \in V_{\mathbb{C}}$ we define $T_p(\mathbf{x}; \mathbf{v})$ by the complex multilinear extension of $D^p u(\mathbf{x})$ on the right-hand side of (2.3). When u admits a holomorphic extension to a neighborhood of $\mathbf{x}$ in $\mathbb{C}^d$, this definition agrees with the order- p coefficient of $\tau \mapsto u(\mathbf{x} + \tau \mathbf{v})$, and the recurrences (2.4)–(2.6) evaluate it on complex tensors. When the directions are complex, the extension is guaranteed by entire activations such as $\sinh(z) = (e^z - e^{-z})/2$; see Remark 3.4. The distinction matters because an arbitrary C^p function on $\mathbb{R}^d$ need not be defined at complex arguments.

3.1. Basic expansion. Let $\mathbf{z} \in \mathbb{C}^{n+1}$ parametrize the active subspace through $\mathbf{v}(\mathbf{z}) := \sum_{j=0}^n z_j e_{i_j}$, and define the *generating polynomial* of u at $\mathbf{x}$ by

$$Q_u(\mathbf{z}) := T_p(\mathbf{x}; \mathbf{v}(\mathbf{z})). \quad (3.2)$$

The symmetric p -linearity of $D^p u(\mathbf{x})$ and the multinomial theorem give

$$\begin{aligned} D^p u(\mathbf{x})[\mathbf{v}(\mathbf{z}), \dots, \mathbf{v}(\mathbf{z})] &= \sum_{|\beta|=p} \binom{p}{\beta} \mathbf{z}^\beta \partial^\beta u(\mathbf{x}) \\ &= \sum_{|\beta|=p} \frac{p!}{\beta!} \mathbf{z}^\beta \partial^\beta u(\mathbf{x}), \end{aligned}$$

where β ranges over multi-indices on the active variables and $\partial^\beta := \partial_{i_0}^{\beta_0} \dots \partial_{i_n}^{\beta_n}$. Dividing by $p!$ and comparing with (2.3) and (3.2) yields

$$Q_u(\mathbf{z}) = \sum_{|\beta|=p} \frac{\partial^\beta u(\mathbf{x})}{\beta!} \mathbf{z}^\beta, \quad (3.3)$$

so Q_u is homogeneous of degree p and its coefficients are the order- p derivatives of u at $\mathbf{x}$. Reading off the coefficient of $\mathbf{z}^a$ in (3.3) and using $a! = \nu!$ from (3.1) gives

$$\partial^\nu u(\mathbf{x}) = \nu! [\mathbf{z}^a] Q_u(\mathbf{z}). \quad (3.4)$$

Extracting one monomial derivative is therefore the extraction of one coefficient of a homogeneous polynomial of degree p .

The evaluations available to us are values of Q_u at individual directions, so we ask for the shortest linear combination of such values that realizes the functional (3.4). A *directional schedule* of length R for ∂^ν over $\mathbb{K}$ is a finite collection $(c_r, \mathbf{v}_r)_{r=1}^R \subset \mathbb{K} \times V_{\mathbb{K}}$ such that

$$\partial^\nu u(\mathbf{x}) = \sum_{r=1}^R c_r T_p(\mathbf{x}; \mathbf{v}_r) \quad (3.5)$$

for every u that is C^p on a neighborhood of $\mathbf{x}$, and, when $\mathbb{K} = \mathbb{C}$, holomorphic there. Because (3.5) is required for all such u , the order- p derivatives $(\partial^\beta u(\mathbf{x}))_{|\beta|=p}$ may be prescribed independently: for any coefficients q_β , the polynomial $\sum_\beta q_\beta \prod_j (y_{i_j} - x_{i_j})^{\beta_j}$ has $\partial^\beta u(\mathbf{x}) = \beta! q_\beta$. Substituting (3.3) into the right-hand side of (3.5) gives

$$\sum_{r=1}^R c_r T_p(\mathbf{x}; \mathbf{v}_r) = \sum_{|\beta|=p} \frac{\partial^\beta u(\mathbf{x})}{\beta!} \left(\sum_{r=1}^R c_r \mathbf{v}_r^\beta \right), \quad (3.6)$$

and comparing (3.6) with the target $\partial^\nu u(\mathbf{x}) = \partial^a u(\mathbf{x})$ coefficient by coefficient shows that (3.5) holds for every admissible u if and only if

$$\sum_{r=1}^R c_r \mathbf{v}_r^\beta = \nu! \delta_{a\beta}, \quad |\beta| = p, \quad (3.7)$$

where $\delta_{a\beta} := 1$ if $\beta = a$ and $\delta_{a\beta} := 0$ otherwise. The design problem is thus to solve the moment conditions (3.7) with as few directions as possible.

Two features of (3.7) decide how short a schedule can be. The conditions number $\binom{p+n}{n}$, so their count is governed by how many coordinates the derivative touches and not by the order alone: a derivative of order six on three coordinates imposes far fewer conditions than one of order six on six coordinates. They are also nonlinear in the unknowns, since each direction enters through the monomials $\mathbf{v}_r^\beta$, which is why the minimum is not read off by counting equations and unknowns. Section 3.2 resolves both by recognizing (3.7) as a polynomial identity whose shortest solution is classical.

3.2. Minimal schedules and monomial Waring rank. Pair a direction $\mathbf{v} \in V_{\mathbb{C}}$, written $\mathbf{v} = \sum_{j=0}^n v_j e_{i_j}$ in active coordinates, with the polynomial variable $\mathbf{z}$ through $\mathbf{v} \cdot \mathbf{z} := \sum_{j=0}^n v_j z_j$. Expanding the p th power of this pairing by the multinomial theorem gives

$$\sum_{r=1}^R c_r (\mathbf{v}_r \cdot \mathbf{z})^p = \sum_{|\beta|=p} \binom{p}{\beta} \left(\sum_{r=1}^R c_r \mathbf{v}_r^\beta \right) \mathbf{z}^\beta, \quad (3.8)$$

whose coefficients are the very quantities constrained by (3.7). The moment conditions are therefore a polynomial identity in disguise, and the shortest schedule is the shortest decomposition of a monomial into pure powers of linear forms. That notion has a classical name.

DEFINITION 3.1 (Waring rank). *Let $\mathbb{K} \in \{\mathbb{R}, \mathbb{C}\}$ and let $F \in \mathbb{K}[\mathbf{z}]_p$ be homogeneous of degree p . The Waring rank $R_{\mathbb{K}}(F)$ is the least R for which there exist $c_r \in \mathbb{K}$ and linear forms $L_r \in \mathbb{K}[\mathbf{z}]_1$ with $F = \sum_{r=1}^R c_r L_r(\mathbf{z})^p$.*

Since $\mathbb{R} \subset \mathbb{C}$, every real decomposition is a complex one, so $R_{\mathbb{C}}(F) \leq R_{\mathbb{R}}(F)$, and enlarging the field can only shorten a schedule. The following theorem is the

main result of the paper. Its three parts identify the design problem with a Waring decomposition, evaluate the resulting minimum, and exhibit a schedule of that length in closed form.

THEOREM 3.2 (Rank-optimal directional schedule). *Let ν be a multi-index of order $p \geq 1$ with active exponents $1 \leq a_0 \leq \dots \leq a_n$ and $m_j := a_j + 1$ as in (3.1), and let $\mathbb{K} \in \{\mathbb{R}, \mathbb{C}\}$.*

- (i) *A collection $(c_r, \mathbf{v}_r)_{r=1}^R \subset \mathbb{K} \times V_{\mathbb{K}}$ is a directional schedule for ∂^ν if and only if*

$$p! \mathbf{z}^a = \sum_{r=1}^R c_r (\mathbf{v}_r \cdot \mathbf{z})^p \quad \text{in } \mathbb{K}[\mathbf{z}]_p. \quad (3.9)$$

- (ii) *The minimum length of a directional schedule for ∂^ν over $\mathbb{K}$ equals $R_{\mathbb{K}}(\mathbf{z}^a)$, and over $\mathbb{C}$ this minimum is*

$$R_{\mathbb{C}}(\mathbf{z}^a) = \prod_{j=1}^n m_j = \frac{\prod_{j=0}^n m_j}{m_0}. \quad (3.10)$$

- (iii) *For $\zeta = (\zeta_1, \dots, \zeta_n) \in U_{m_1} \times \dots \times U_{m_n}$, put*

$$\mathbf{v}_\zeta := e_{i_0} + \sum_{j=1}^n \zeta_j e_{i_j}, \quad c_\zeta := \frac{\nu!}{\prod_{j=1}^n m_j} \prod_{j=1}^n \zeta_j. \quad (3.11)$$

Then $(c_\zeta, \mathbf{v}_\zeta)_\zeta$ is a directional schedule for ∂^ν of length $\prod_{j=1}^n m_j$, and by part (ii) no complex schedule is shorter.

Proof.

(i) The moment conditions are a polynomial identity. By (3.7), the collection is a directional schedule if and only if $\sum_r c_r \mathbf{v}_r^\beta = \nu! \delta_{a\beta}$ for every β with $|\beta| = p$. By (3.8), the identity (3.9) holds if and only if $\binom{p}{\beta} \sum_r c_r \mathbf{v}_r^\beta = p! \delta_{a\beta}$ for every such β . Dividing the second family of conditions by $\binom{p}{\beta}$ turns it into $\sum_r c_r \mathbf{v}_r^\beta = \beta! \delta_{a\beta}$, and $\beta! \delta_{a\beta} = \nu! \delta_{a\beta}$ by (3.1). The two families coincide, which proves part (i).

(ii) The minimum length is a Waring rank. Absorbing the nonzero constant $p!$ into the coefficients c_r shows that a decomposition of $p! \mathbf{z}^a$ into R pure p th powers exists if and only if one of $\mathbf{z}^a$ does, so $R_{\mathbb{K}}(p! \mathbf{z}^a) = R_{\mathbb{K}}(\mathbf{z}^a)$. Combining this with part (i) identifies the minimum schedule length over $\mathbb{K}$ with $R_{\mathbb{K}}(\mathbf{z}^a)$. The value (3.10) of the complex rank of a monomial is the theorem of Carlini, Catalisano, and Geramita [2]; Appendix A.3 recalls the apolarity argument behind its lower bound.

(iii) The roots-of-unity schedule attains the minimum. We use the conventions that an empty product equals 1 and that a Cartesian product over an empty index set has one element, the empty tuple; these cover the case $n = 0$, in which (3.11) produces the single direction e_{i_0} with coefficient $\nu!$. By part (i) it suffices to prove the identity $\sum_\zeta c_\zeta (\mathbf{v}_\zeta \cdot \mathbf{z})^p = p! \mathbf{z}^a$. Fix β with $|\beta| = p$ and let C_β denote the coefficient of $\mathbf{z}^\beta$ on the left-hand side. Inserting (3.11) and applying the multinomial theorem gives

$$\begin{aligned} C_\beta &= \frac{\nu!}{\prod_{j=1}^n m_j} \binom{p}{\beta} \sum_{\zeta_1 \in U_{m_1}} \dots \sum_{\zeta_n \in U_{m_n}} \prod_{j=1}^n \zeta_j^{\beta_j+1} \\ &= \frac{\nu!}{\prod_{j=1}^n m_j} \binom{p}{\beta} \prod_{j=1}^n \left(\sum_{\zeta_j \in U_{m_j}} \zeta_j^{\beta_j+1} \right), \end{aligned} \quad (3.12)$$

where the second equality separates the independent sums. By Lemma A.1, the j th factor of the product in (3.12) vanishes unless $m_j \mid \beta_j + 1$, and $m_j = a_j + 1$ turns this divisibility into the congruence $\beta_j \equiv a_j \pmod{m_j}$. Suppose $C_\beta \neq 0$, and write $\beta_j = a_j + t_j m_j$ with $t_j \in \mathbb{N}_0$ for $j = 1, \dots, n$. The degree constraint $|\beta| = p = \sum_{j=0}^n a_j$ then forces $\beta_0 = a_0 - \sum_{j=1}^n t_j m_j$. If some t_j were positive, then $m_j = a_j + 1 \geq a_0 + 1$ by the ordering of the exponents, whence $\beta_0 \leq a_0 - m_j < 0$, contradicting $\beta_0 \geq 0$. Every t_j therefore vanishes, so $\beta_j = a_j$ for $j \geq 1$, and the degree constraint gives $\beta_0 = a_0$ as well. Hence $C_\beta = 0$ for every $\beta \neq a$. For $\beta = a$ each inner sum in (3.12) equals m_j , the product cancels the prefactor, and

$$C_a = \nu! \binom{p}{a} = \left(\prod_{j=0}^n a_j! \right) \frac{p!}{\prod_{j=0}^n a_j!} = p!.$$

The left-hand side of the identity therefore equals $p! z^a$, and the schedule has $|U_{m_1} \times \dots \times U_{m_n}| = \prod_{j=1}^n m_j$ directions. $\square$

Theorem 3.2 converts a question about derivative evaluation into a rank computation whose answer is the product (3.10). The minimizing directions and coefficients in (3.11) are roots of unity scaled by $\nu! / R_{\mathbb{C}}(z^a)$, so their magnitudes are known in advance. Three consequences deserve separate mention.

REMARK 3.3 (The smallest instance). *For $\partial^\nu = \partial_1 \partial_2$ we have $n = 1$, $a_0 = a_1 = 1$, $m_1 = 2$, and $U_2 = \{+1, -1\}$, so the roots of unity are real and (3.11) returns the two pairs $(\frac{1}{2}, e_1 + e_2)$ and $(-\frac{1}{2}, e_1 - e_2)$, that is, the elementary identity*

$$\partial_1 \partial_2 u(\mathbf{x}) = \frac{1}{2} T_p(\mathbf{x}; e_1 + e_2) - \frac{1}{2} T_p(\mathbf{x}; e_1 - e_2), \quad (3.13)$$

in which the pure second derivatives cancel between the two probes. Theorem 3.2 adds to (3.13) the information that two evaluations are necessary as well as sufficient. Complex directions appear once two or more coordinates are active and some exponent is at least two. A pure power $\partial_1^p u$ uses one real direction, and a square-free monomial uses only the real pair $U_2 = \{\pm 1\}$.

REMARK 3.4 (Entire activations are essential). *The directions (3.11) are complex whenever some factor U_{m_j} with $j \geq 1$ has $m_j \geq 3$, so the composition is then evaluated at complex arguments and its activations must admit a holomorphic extension there. The functions $\sinh$ and $\cosh$ are entire, and the recurrence (2.6) is an algebraic identity that holds verbatim on complex tensors. By contrast $\tanh$ is meromorphic, with poles at $i\pi/2 + i k\pi$ for $k \in \mathbb{Z}$, so its jet is defined only away from those points and the admissible directions depend on the activation input. The collocation measurements of Section 4 use $\sinh$ at every order, including the order-two residuals whose directions remain real.*

REMARK 3.5 (Scope of the rank reduction). *Theorem 3.2 concerns one monomial derivative at a time. For the square-free multi-index $a = (1, \dots, 1)$ of order p the rank (3.10) equals 2^{p-1} , so the schedule is exponentially long; since the theorem gives a minimum, this is a statement about the problem and not about the construction, and no schedule of the form (3.5) can be shorter. A differential operator that is a linear combination of monomial derivatives is evaluated by applying the schedule to each monomial. The stochastic Taylor derivative estimator [14] treats the strictly broader problem of an arbitrary contraction of the order- p derivative tensor, with unbiased estimates whose variance is controlled by a sampling budget; the two approaches are complementary, and the schedule of Theorem 3.2 is the deterministic minimum for one fixed multi-index.*

Table 3.1 evaluates (3.10) for selected active-exponent patterns up to $p = 6$. Because every entry is a minimum, the table reads as a statement about which derivatives are cheap to extract from directional evaluations and which are not: a pure power needs one direction at any order, a derivative that repeats every index stays modest as the order grows, and a square-free derivative is intrinsically exponential. The explicit schedules for the patterns that occur in Section 4 are listed in Appendix A.2.

TABLE 3.1

The minimum number of directional evaluations $R_{\mathbb{C}}(\mathbf{z}^a)$ of (3.10), for selected active-exponent patterns up to $p = 6$. Repeating an index keeps the count low; the square-free pattern (1^p) is exponential in the order.

pattern	example multi-index	$R_{\mathbb{C}}$
(p)	$\partial_1^p u$	1
$(2, 1)$	u_{112}	3
$(1, 1, 1)$	u_{123}	4
$(3, 1)$	u_{1112}	4
$(2, 2)$	u_{1122}	3
$(2, 1, 1)$	u_{1123}	6
$(1, 1, 1, 1)$	u_{1234}	8
$(4, 2)$	u_{111122}	5
$(2, 2, 2)$	u_{112233}	9
(1^6)	u_{123456}	32

3.3. Algorithm, gradients, and cost. Algorithm 1 assembles the pieces. The schedule (3.11) is built once, its directions are stacked into a single batch, and one jet pass through the composition returns the order- p coefficient for every direction at once; the reported value is then the coefficient-weighted sum. The schedule tensors depend on ν , the device, and the scalar type, but not on the parameters of the composition. The reference implementation builds them at call entry and includes that work in the measured throughput of Section 4, while a caller that reuses one multi-index may cache them. For a real-valued composition the conjugate pairing of the roots of unity makes the returned value real in exact arithmetic; the implementation keeps the complex tensor and takes its real part only when the target PDE is real-valued.

Algorithm 1 must be differentiable in θ , since its output enters the residual term of (2.2). Differentiating the activation recurrences (2.5)–(2.6) by reverse-mode automatic differentiation would rebuild inside the backward pass the nesting that the jet was introduced to avoid. We therefore register each activation jet as a single primitive with the closed-form adjoint

$$g_x^{(m)} = \sum_{j=0}^{p-m} \overline{q_j} g_y^{(m+j)}, \quad q(\tau) := \phi'(x(\tau)), \quad (3.14)$$

in which q_j is the j th Taylor coefficient of q and $g^{(m)}$ denotes the cotangent of the order- m coefficient. Equation (3.14) follows from the chain rule applied to the jet of ϕ and uses the convention of a real-valued loss for complex parameters, so that the conjugation is the Hermitian adjoint of multiplication by q_j . The backward pass thus costs one more triangular convolution per activation and never differentiates through the recurrence itself.

The cost of Algorithm 1 separates into a per-direction and a per-schedule factor. For one direction, the jet stores $p + 1$ coefficient tensors per activation, an affine layer

Algorithm 1 Single-monomial $\partial^\nu u(\mathbf{x})$ from a rank-optimal schedule and one batched Taylor jet.

Require: scalar network u built from supported analytic primitives; point $\mathbf{x} \in \mathbb{R}^d$; multi-index ν of order p

Ensure: the value $\partial^\nu u(\mathbf{x})$

- 1: Relabel the active coordinates of ν as $i_0, \dots, i_n$ with exponents $1 \leq a_0 \leq \dots \leq a_n$.
 - 2: Set $m_j := a_j + 1$ for $j = 1, \dots, n$ and $\mathcal{I} := U_{m_1} \times \dots \times U_{m_n}$.
 - 3: Initialize $\mathcal{S} \leftarrow \emptyset$.
 - 4: **for** $\zeta = (\zeta_1, \dots, \zeta_n) \in \mathcal{I}$ **do**
 - 5: Set $\mathbf{v}_\zeta \leftarrow e_{i_0} + \sum_{j=1}^n \zeta_j e_{i_j}$ and $c_\zeta \leftarrow \nu! (\prod_{j=1}^n m_j)^{-1} \prod_{j=1}^n \zeta_j$.
 - 6: Update $\mathcal{S} \leftarrow \mathcal{S} \cup \{(c_\zeta, \mathbf{v}_\zeta)\}$.
 - 7: **end for**
 - 8: Stack the directions of $\mathcal{S}$ into $V_\mathcal{S} \in \mathbb{C}^{|\mathcal{S}| \times d}$.
 - 9: Form the batched input jet $\mathbf{J}_0 \leftarrow (\mathbf{x} \text{ replicated}, V_\mathcal{S}, 0, \dots, 0)$.
 - 10: Propagate $\mathbf{J}_0$ through u by (2.4)–(2.6) to obtain $\mathbf{J}_L$.
 - 11: Set $t \leftarrow [\mathbf{J}_L]_p$, one order- p coefficient per direction.
 - 12: **return** $\sum_{r=1}^{|\mathcal{S}|} c_r t_r$.
-

applies W to each of them, and the activation recurrences require $\sum_{k=1}^p O(k) = O(p^2)$ elementwise products per scalar. The total work and the directional storage then scale linearly in the schedule length $R_\mathbb{C}(\mathbf{z}^a)$, which by Theorem 3.2 is as small as an exact directional formula can be. The rank reduction is an exact statement about that factor and not about wall-clock time, which also depends on complex arithmetic, kernel launch overhead, memory traffic, and whether parameter gradients are required; for small widths or small batches these terms can dominate. Section 4 therefore reports measured throughput alongside accuracy.

The schedule coefficients are known in closed form. The remaining error is floating-point round-off from the Taylor recurrences, from complex arithmetic, and from the final summation over directions. Section 4.1 reports double-precision agreement through order six.

4. Numerical experiments. This section checks two properties of the schedule: agreement with nested automatic differentiation in double precision, and the time and memory of one residual evaluation, including the reverse-mode pass. The first comparison uses no PDE. The second repeats the same evaluation inside an otherwise identical collocation loop and records the time and memory of one step. Every quantity below is a property of the derivative method on a fixed composition, residual, and optimizer.

4.1. Verification of the derivative backend. All checks use double precision and are executable tests in the repository. The roots-of-unity coefficient filter of Lemma A.1 is evaluated against every degree-four monomial in five variables for the active-exponent pattern $(2, 1, 1)$, with a maximum admissible absolute coefficient error of 10^{-12} . The length of the schedule that the implementation builds is checked against the closed-form rank (3.10) for every pattern of Table 3.1.

The Waring jet of Algorithm 1 is then compared with nested automatic differentiation on real-parameter networks using the tanh, logistic, sin, and sinh activations, for the patterns $(2, 2)$, $(1, 1, 1)$, and (3) . Values must agree to relative tolerance 2×10^{-10}

and absolute tolerance 2×10^{-11} , and the imaginary remainder of a real target must not exceed 2×10^{-11} . Parameter gradients of a squared derivative loss are compared at tolerances 2×10^{-9} relative and 2×10^{-11} absolute, on a real tanh network and on a complex sinh network; additional analytic checks cover the orders 1, 2, 3, 4, and 6 for a scalar complex sinh network. Every check passes at the tolerance stated for it. These tests establish that the implementation computes what Theorem 3.2 predicts.

4.2. A controlled comparison of derivative backends. The remaining study isolates the derivative backend. Protocol `jsc.v3` runs two lines that share the network, the loss, the optimizer, the collocation points, the evaluation rule, the seeds, and the wall-clock budget, and differ in one respect: how each monomial derivative of the residual is evaluated. The first line uses Algorithm 1, that is, the rank-optimal complex schedule evaluated by one batched Taylor jet, and is labeled *jet* below. The second evaluates the same derivative of the same network by nested reverse-mode automatic differentiation, with the graph retained after every coordinate derivative, and is labeled *nested*. Both lines are the same complex-parameter sinh network with four hidden layers of width $H = 128$, hence the same 100,098 trainable real degrees of freedom, the same first-layer scale, and the same `complex128` arithmetic. Both use Adam with initial learning rate 10^{-3} and cosine decay, 4096 fixed interior points, 512 fixed boundary points, paired per-seed collocation, the same per-setting boundary weights, and a common held-out set of 8192 points. Each line is run for three seeds with a budget of 1000 seconds per run.

A *setting* is one problem together with one value of its sweep parameter, and the protocol covers nine: the two-dimensional polyharmonic problem at orders $2m \in \{2, 4, 6\}$, the radial chirp at $a \in \{1, 2, 3\}$, and the time-harmonic Maxwell problem at $a \in \{2, 4, 6\}$, all defined as in Section 4.3. All 54 runs were produced by one source commit on one NVIDIA H20 GPU with PyTorch 2.5.1 and CUDA 12.1, so unlike a comparison assembled from separate campaigns, timings here are comparable across settings as well as within one. Cost is the measured time per optimizer step together with the peak memory allocated by a run, and accuracy is the held-out relative L^2 error

$$\mathcal{E}_{L^2} := \frac{\|u_\theta - u_\star\|_{L^2(\Omega)}}{\|u_\star\|_{L^2(\Omega)}}, \quad (4.1)$$

evaluated on that held-out set against the exact solution $u_\star$. A setting enters the tables only after automated checks confirm all six method-seed rows, finite metrics, nonempty monotone histories, and the protocol metadata.

4.3. Benchmark problems. The three families place different demands on the derivative. For the polyharmonic problem on $(-1, 1)^2$ we solve

$$\Delta^m u = (-2\pi^2)^m u_\star \quad \text{in } (-1, 1)^2, \quad u_\star(\mathbf{x}) = \sin(\pi x_1) \sin(\pi x_2), \quad (4.2)$$

at the orders $2m \in \{2, 4, 6\}$, with the Navier boundary targets $u = u_\star$ and $\Delta^j u = (-2\pi^2)^j u_\star$ for $j = 1, \dots, m-1$. Those traces vanish on $\partial(-1, 1)^2$ because $u_\star$ does. This family is the reason the derivative order appears as a variable: expanding Δ^m in two dimensions produces the patterns $(2m)$ and, for $m \geq 2$, the mixed patterns whose schedules Table 3.1 lists. At $2m = 4$ the mixed term is $\partial_1^2 \partial_2^2$, which (3.11) evaluates with three directions; at $2m = 6$ the mixed terms are $\partial_1^4 \partial_2^2$ and $\partial_1^2 \partial_2^4$, which take five each.

The other two families sit at order two and sweep a frequency instead. On $(-1, 1)^2$ the radial chirp is

$$-\Delta u + u = f, \quad u_*(x, y) = \sin\left(\frac{a\pi}{2}(x^2 + y^2)\right), \quad (4.3)$$

with Dirichlet data from u_* , $a \in \{1, 2, 3\}$, and the manufactured source $f = (a\pi)^2 r^2 \sin \phi - 2a\pi \cos \phi + \sin \phi$, where $\phi := a\pi(x^2 + y^2)/2$ and $r^2 := x^2 + y^2$. Its local radial wavenumber $|\nabla \phi| = a\pi r$ varies over the domain, so the target is not a separable Fourier mode. For a two-dimensional transverse-magnetic mode, Maxwell's equations reduce to a scalar complex Helmholtz equation for the out-of-plane field E ,

$$\Delta E + \kappa^2 E = f, \quad \kappa^2 = (a\pi)^2(1 + i\beta), \quad E_* = \exp(ia\pi(x + y)), \quad (4.4)$$

with $\beta = 0.2$, exact Dirichlet data, $a \in \{2, 4, 6\}$, and $f = [-2(a\pi)^2 + \kappa^2]E_*$. The Maxwell family has a complex-valued solution. The measurements already propagate complex tensors through a complex-parameter sinh composition, so a complex residual uses the same jet as the real ones. Since both families are of order two, every monomial term of their residuals is a pure power, which the schedule evaluates with a single real direction. They therefore isolate the cost of the Taylor jet on one direction.

TABLE 4.1

Cost of one optimizer step under the jsc_v3 protocol, which changes only the derivative backend. Times are the three-seed mean of milliseconds per step and memory is the peak allocation of a run; the ratio columns give nested automatic differentiation divided by the Taylor jet, so larger is better for the jet.

setting	p	ms per step			peak memory (MiB)		
		jet	nested	ratio	jet	nested	ratio
Polyharmonic, $2m = 2$	2	62.3	91.1	1.46	509	789	1.55
Polyharmonic, $2m = 4$	4	251.6	565.9	2.25	1759	5400	3.07
Polyharmonic, $2m = 6$	6	854.2	3496.1	4.09	5728	36976	6.46
Radial chirp, $a = 1$	2	72.8	101.7	1.40	548	852	1.55
Radial chirp, $a = 2$	2	73.4	101.9	1.39	548	852	1.55
Radial chirp, $a = 3$	2	73.4	102.1	1.39	548	852	1.55
Maxwell, $a = 2$	2	73.3	102.1	1.39	573	853	1.49
Maxwell, $a = 4$	2	73.2	101.9	1.39	573	853	1.49
Maxwell, $a = 6$	2	72.9	101.8	1.40	573	853	1.49

4.4. Cost of one derivative evaluation. Table 4.1 and Figure 4.1 report the cost of one optimizer step. At order two the jet is faster by a factor between 1.39 and 1.46 across all seven order-two settings, and it allocates between 1.49 and 1.55 times less memory. At order two every residual term is a pure power, so the schedule uses one direction and the saving is that of Taylor-mode evaluation: nested differentiation retains a graph across two coordinate derivatives. The saving then grows with the order, exactly where the rank reduction begins to apply: at order four the step time falls by a factor of 2.25 and the peak memory by 3.07, and at order six by 4.09 and 6.46. In absolute terms, one step of the sixth-order problem costs 854 milliseconds and 5.6 gibibytes with the jet, against 3.5 seconds and 36.1 gibibytes with nested differentiation.

Two mechanisms compound in that growth, and the design of the study separates them. The jet replaces a graph of depth proportional to the derivative order and the network depth by $p + 1$ coefficient tensors per activation, which is what the memory




FIG. 4.1. Cost of one optimizer step against the derivative order, on the polyharmonic family, for the two backends of Section 4.2. Panels (a) and (b) show the three-seed mean step time and the peak memory of a run; panel (c) shows the ratio of nested automatic differentiation to the Taylor jet for both quantities.

column measures; this part is available at every order and is what the order-two settings show. The schedule then fixes how many times that jet must be run, at the minimum of Table 3.1: one direction for each pure power at any order, three for the mixed term at order four, five for each mixed term at order six. Nested differentiation has no comparable lever, since it must traverse the retained graph once per coordinate derivative of every term, which is why the ratio in panel (c) of Figure 4.1 widens with the order. Both counts are exact. Table 4.1 shows that on this hardware the reduction survives complex arithmetic, kernel launches, and the backward pass.

TABLE 4.2

What the saved time buys under the same protocol: optimizer steps completed inside the common 1000-second budget, as a three-seed mean, and the held-out relative L^2 error reached, as a three-seed mean $\pm$ sample standard deviation. The lower error in each row is set in bold.

setting	steps completed		relative L^2 error	
	jet	nested	jet	nested
Polyharmonic, $2m = 2$	16058	10973	$3.25 \pm 1.15 \times 10^{-4}$	$5.42 \pm 0.88 \times 10^{-4}$
Polyharmonic, $2m = 4$	3975	1767	$6.55 \pm 0.71 \times 10^{-4}$	$1.43 \pm 0.10 \times 10^{-3}$
Polyharmonic, $2m = 6$	1171	287	$3.36 \pm 1.81 \times 10^{-2}$	$9.18 \pm 3.62 \times 10^{-2}$
Radial chirp, $a = 1$	13730	9835	$3.84 \pm 3.68 \times 10^{-4}$	$7.05 \pm 3.55 \times 10^{-4}$
Radial chirp, $a = 2$	13617	9816	$8.93 \pm 10.97 \times 10^{-4}$	$8.36 \pm 4.56 \times 10^{-4}$
Radial chirp, $a = 3$	13629	9796	$1.86 \pm 0.23 \times 10^{-4}$	$3.38 \pm 1.25 \times 10^{-4}$
Maxwell, $a = 2$	13639	9795	$1.24 \pm 0.44 \times 10^{-3}$	$1.61 \pm 0.72 \times 10^{-3}$
Maxwell, $a = 4$	13665	9814	$1.24 \pm 0.40 \times 10^{-3}$	$1.48 \pm 0.67 \times 10^{-3}$
Maxwell, $a = 6$	13712	9825	$9.27 \pm 1.52 \times 10^{-3}$	$9.42 \pm 1.49 \times 10^{-3}$

4.5. What the saving buys. A cheaper evaluation is useful if more evaluations fit in the same wall-clock budget. Table 4.2 therefore records how many optimizer steps each line completes inside the common budget and the accuracy it reaches, and Figure 4.2 shows the trajectories. The step counts follow the cost directly: at order six the jet completes 1171 steps against 287, and at order four 3975 against 1767. The




FIG. 4.2. Held-out relative L^2 error against training time for the two backends, on one representative setting of each family. Curves are the three-seed median and the shaded band is the seedwise range. Each panel is framed on the median curves.

accuracy follows the step counts where the problem is hard enough for the extra steps to matter. At order four the jet reaches 6.55×10^{-4} against 1.43×10^{-3} , and at order six 3.36×10^{-2} against 9.18×10^{-2} , in both cases with the seed spread of Table 4.2 well inside the gap. Across the nine settings the jet has the lower three-seed mean error in eight; the exception is the chirp at $a = 2$, where the two lines are within one seed standard deviation of each other.

At order two the extra steps buy little: the two backends track each other in panels (b) and (c) of Figure 4.2, and in five of the six frequency-swept settings the two means differ by less than one seed standard deviation. A cheaper derivative is not by itself a better optimizer, and at low order the derivative is not what limits the run. For a prescribed monomial derivative the schedule is algebraically exact, and its cost advantage grows with the derivative order, in wall time and in memory, with the model, the loss, and the optimizer held fixed.

5. Conclusion. Evaluating one prescribed high-order derivative exactly is a question about how few directions suffice, and we have shown that it is the Waring problem for the associated monomial. The identification is an equivalence rather than a bound: a directional formula of a given length exists precisely when the monomial decomposes into that many powers of linear forms. The rank formula of Carlini, Catalisano, and Geramita then supplies the complex minimum (3.10), and a roots-of-unity construction attains it in closed form. Because the rank is a minimum, a shorter linear combination of order- p directional Taylor coefficients is impossible, and the elementary two-point identity for $\partial_1 \partial_2$ is the smallest instance.

Combined with a batched Taylor jet through an entire activation and a closed-form vector-Jacobian product, the schedule evaluates the derivative exactly and remains differentiable in the parameters of the composition. Double-precision tests reproduce nested automatic differentiation for values and for parameter gradients through order six. Measured against nested differentiation on collocation residuals where nothing else varies, the schedule reduces the time of one evaluation by factors of about 1.4, 2.3, and 4.1 at derivative orders two, four, and six, and its peak mem-

ory by about 1.5, 3.1, and 6.5; on the higher-order problems the evaluations thus gained lower the error reached within a fixed budget. The two mechanisms behind that growth are separable in the measurements: Taylor-mode evaluation removes the nested graph at every order, and the rank reduction shortens the schedule wherever an index repeats.

The schedule is shortest for repeated-index multi-indices and exponentially long for square-free ones: Theorem 3.2 shows that this length is the minimum. At low derivative order a cheaper derivative also buys little in a complete run, since the derivative is not what limits it. Natural continuations are a rank-aware expansion of an operator that reuses directions across its terms, a conditioning analysis of the schedule at high order and in reduced precision, and the extension of the same apolarity argument to derivative tensors contracted against a fixed structured family.

Appendix A. Schedules, filters, and the apolar bound.

This appendix collects the roots-of-unity filter used in the proof of Theorem 3.2, the explicit schedules for the derivative patterns that occur in Section 4, and the apolarity argument behind the rank formula (3.10). The formulas of Appendix A.2 are those the implementation evaluates.

A.1. Roots-of-unity filter. LEMMA A.1. *For every integer $m \geq 1$ and every integer k ,*

$$\sum_{\zeta \in U_m} \zeta^k = m \mathbf{1}_{m|k}. \quad (\text{A.1})$$

Proof. Let $\omega := \exp(2\pi i/m)$, so that every element of U_m is ω^ℓ for exactly one $\ell \in \{0, \dots, m-1\}$ and

$$\sum_{\zeta \in U_m} \zeta^k = \sum_{\ell=0}^{m-1} (\omega^k)^\ell.$$

If $m \mid k$, then $\omega^k = 1$ and the sum equals m . If $m \nmid k$, then $\omega^k \neq 1$, and the finite geometric series gives

$$\sum_{\ell=0}^{m-1} (\omega^k)^\ell = \frac{1 - \omega^{km}}{1 - \omega^k} = 0,$$

since $\omega^m = 1$. □

A.2. Explicit low-order schedules. The mixed terms of the polyharmonic residual (4.2) and the pure powers of the chirp and Maxwell residuals use the following instances of (3.11).

For a pure power $\partial^\nu = \partial_1^p$ we have $n = 0$, the schedule has one direction, and T_p is evaluated once,

$$\partial_1^p u(\mathbf{x}) = p! T_p(\mathbf{x}; e_1). \quad (\text{A.2})$$

For the fourth-order mixed term $\partial^\nu = \partial_1^2 \partial_2^2$ we have $n = 1$, $a = (2, 2)$, and $m_1 = 3$; with $\omega := \exp(2\pi i/3)$ and $U_3 = \{1, \omega, \omega^2\}$, formula (3.11) gives the three directions

$$\partial_1^2 \partial_2^2 u(\mathbf{x}) = \frac{4}{3} \sum_{\zeta \in U_3} \zeta T_p(\mathbf{x}; e_1 + \zeta e_2). \quad (\text{A.3})$$

For the sixth-order mixed term $\partial^\nu = \partial_1^2 \partial_2^4$ we have $n = 1$, $a = (2, 4)$, and $m_1 = 5$; with U_5 the fifth roots of unity, (3.11) takes the coordinate of smaller exponent as base and gives the five directions

$$\partial_1^2 \partial_2^4 u(\mathbf{x}) = \frac{48}{5} \sum_{\zeta \in U_5} \zeta T_p(\mathbf{x}; e_1 + \zeta e_2). \quad (\text{A.4})$$

The companion term $\partial_1^4 \partial_2^2$ is the same formula with the two coordinates exchanged. The equivalent polynomial identity (3.9) for (A.4) reads

$$6! z_0^2 z_1^4 = \frac{48}{5} \sum_{\zeta \in U_5} \zeta (z_0 + \zeta z_1)^6, \quad (\text{A.5})$$

and it can be checked directly: the coefficient of $z_0^{\beta_0} z_1^{\beta_1}$ on the right-hand side contains the factor $\sum_{\zeta \in U_5} \zeta^{1+\beta_1}$, which by Lemma A.1 is nonzero only when $\beta_1 \equiv 4 \pmod{5}$. Together with $\beta_0 + \beta_1 = 6$ this forces $\beta = (2, 4)$, and the surviving constant is $\frac{48}{5} \cdot \frac{6!}{2!4!} \cdot 5 = 720 = 6!$. The three-variable pattern $(2, 2, 2)$ of Table 3.1 is the same construction with two cube-root factors and nine directions.

A.3. Apolar form of the rank bound. The lower bound in (3.10) comes from apolarity, which we recall for completeness. For a homogeneous $F \in \mathbb{C}[\mathbf{z}]_p$, let $F^\perp$ denote its apolar ideal, that is, the ideal of differential operators annihilating F , and call a finite subscheme $X \subset \mathbb{P}^n$ apolar to F when $I_X \subset F^\perp$. The Waring rank then admits the characterization

$$R_{\mathbb{C}}(F) = \min\{\deg X : X \subset \mathbb{P}^n \text{ finite, reduced, and apolar to } F\}. \quad (\text{A.6})$$

For the monomial $F = \mathbf{z}^a$ the apolar ideal is the complete intersection

$$(\mathbf{z}^a)^\perp = (\partial_0^{a_0+1}, \dots, \partial_n^{a_n+1}), \quad (\text{A.7})$$

and the Hilbert function of the quotient $\mathbb{C}[\partial]/(\mathbf{z}^a)^\perp$ bounds the degree of an apolar subscheme from below by the product $\prod_{j=1}^n m_j$ of (3.10); see [9, 11] for the general theory and [2] for the monomial case. The construction (3.11) exhibits a reduced apolar subscheme of exactly that degree, so the bound is attained.

REFERENCES

- [1] JESSE BETTENCOURT, MATTHEW J. JOHNSON, AND DAVID DUVENAUD, *Taylor-mode automatic differentiation for higher-order derivatives in JAX*, in NeurIPS Workshop on Program Transformations for Machine Learning, 2019.
- [2] ENRICO CARLINI, MARIA VIRGINIA CATALISANO, AND ANTHONY V. GERAMITA, *The solution to the Waring problem for monomials and the sum of coprime monomials*, Journal of Algebra, 370 (2012), pp. 5–14.
- [3] PAO-HSIUNG CHIU, JIAN CHENG WONG, CHINCHUN OOI, MY HA DAO, AND YEW-SOON ONG, *CAN-PINN: A fast physics-informed neural network based on coupled-automatic-numerical differentiation method*, Computer Methods in Applied Mechanics and Engineering, 395 (2022), p. 114909.
- [4] NICHOLAS GAO, JONAS KÖHLER, AND ADAM FOSTER, *folx: Forward Laplacian for JAX*, 2023. Software.
- [5] ANDREAS GRIEWANK AND ANDREA WALTHER, *Evaluating Derivatives: Principles and Techniques of Algorithmic Differentiation*, Society for Industrial and Applied Mathematics, Philadelphia, PA, 2nd ed., 2008.

- [6] DI HE, SHANDA LI, WENLEI SHI, XIAOTIAN GAO, JIA ZHANG, JIANG BIAN, LIWEI WANG, AND TIE-YAN LIU, *Learning physics-informed neural networks without stacked back-propagation*, in Proceedings of the 26th International Conference on Artificial Intelligence and Statistics, vol. 206 of Proceedings of Machine Learning Research, PMLR, 2023, pp. 3034–3047.
- [7] ZHEYUAN HU, ZEKUN SHI, GEORGE EM KARNIADAKIS, AND KENJI KAWAGUCHI, *Hutchinson trace estimation for high-dimensional and high-order physics-informed neural networks*, Computer Methods in Applied Mechanics and Engineering, 424 (2024), p. 116883.
- [8] ZHEYUAN HU, KHEMRAJ SHUKLA, GEORGE EM KARNIADAKIS, AND KENJI KAWAGUCHI, *Tackling the curse of dimensionality with physics-informed neural networks*, Neural Networks, 176 (2024), p. 106369.
- [9] ANTHONY IARROBINO AND VASSIL KANEV, *Power Sums, Gorenstein Algebras, and Determinantal Loci*, vol. 1721 of Lecture Notes in Mathematics, Springer, Berlin, 1999.
- [10] GEORGE EM KARNIADAKIS, IOANNIS G. KEVREKIDIS, LU LU, PARIS PERDIKARIS, SIFAN WANG, AND LIU YANG, *Physics-informed machine learning*, Nature Reviews Physics, 3 (2021), pp. 422–440.
- [11] JOSEPH M. LANDSBERG, *Tensors: Geometry and Applications*, vol. 128 of Graduate Studies in Mathematics, American Mathematical Society, Providence, RI, 2012.
- [12] RUICHEN LI, HAOTIAN YE, DU JIANG, XUELAN WEN, CHUWEI WANG, ZHE LI, XIANG LI, DI HE, JI CHEN, WEILUO REN, AND LIWEI WANG, *A computational framework for neural network-based variational Monte Carlo with Forward Laplacian*, Nature Machine Intelligence, 6 (2024), pp. 209–219.
- [13] MAZIAR RAISSI, PARIS PERDIKARIS, AND GEORGE E. KARNIADAKIS, *Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations*, Journal of Computational Physics, 378 (2019), pp. 686–707.
- [14] ZEKUN SHI, ZHEYUAN HU, MIN LIN, AND KENJI KAWAGUCHI, *Stochastic Taylor derivative estimator: Efficient amortization for arbitrary differential operators*, Advances in Neural Information Processing Systems (NeurIPS), 37 (2024), pp. 122316–122353.